

Nicholas Prezioso

njp107 Michael

Serpico mvs71

CS214 Asst2 Readme

We wrote two programs that do Run Length Encoding using multiple threads and processes respectively. In order to be more efficient, this process was only done to characters repeated at least 3 times. The programs split the input file into several parts, each of which is simultaneously compressed and output into its own file. Both versions can be run using two command line arguments (the path to the file to be compressed and the number of parts to split the compression into). In both programs, an error is printed if the user tries to split the input file into a number of parts greater than the number of alphabetic characters in the input file. When the input file is read, all non-alphabetic characters are completely ignored. This means that aa\$aa becomes 4a. Each of the given arguments in both programs is checked for validity, and error messages are printed when applicable. In the process version, we decided to have the parent process spawn all of the children and then wait for them all rather than having each process spawn a single child and wait for it.